



Service Mesh & Security

On Container-based Schedulers

1. Digital Transformation: IT World is changing



<https://www.slideshare.net/asotobu/sail-in-the-cloud>

1. Digital Transformation: New challenges

DevOps Challenges for Multiple Containers

- How to scale?
- How to avoid port conflicts?
- How to manage them on multiple hosts?
- What happens if a host has trouble?
- How to keep them running?
- How to update them?
- Where are my containers?




10 @alexstob

RED HAT
DEVELOPERS

<https://www.slideshare.net/asotobu/sail-in-the-cloud>

2. Choosing the Container-based Scheduler




3. Kubernetes




3.1. Kubernetes Architecture

<https://platform9.com/blog/kubernetes-docker-swarm-compared>






3.2. Kubernetes effect: Primitives (1/2)



Concern	Java & JVM	Kubernetes
Behaviour encapsulation	Class	Container Image
Behaviour instance	Object	Container
Unit of reuse	.jar	Container Image
Deployment unit	.jar/.war/.ear	Pod
Buildtime/Runtime isolation	Module, Package, Class	Container Image, Namespace
Initialization preconditions	Constructor	Init-container
Post initialization	init-method	PostStart
Pre destroy	destroy-method	PreStop
Cleanup procedure	finalize(), ShutdownHook	Defer-container*
Asynchronous & Parallel execution	ThreadPoolExecutor, ForkJoinPool	Job
Periodic task	Timer, ScheduledExecutorService	CronJob
Background task	Daemon Thread	DaemonSet
Configuration management	System.getenv(), Properties	ConfigMap, Secret



3.3. Kubernetes effect: Primitives (2/2)










<https://www.infoq.com/articles/kubernetes-effect>

<https://www.slideshare.net/asotobu/sail-in-the-cloud>


3.4. Kubernetes effect: Principles



3.5. Kubernetes effect: Patterns



3.6. Kubernetes effect: Best Practices




Practices & Techniques



In addition to the principles and patterns, creating good containerized applications requires familiarity with other container-related best practices and techniques. Principles and patterns are abstract, fundamental ideas that change less often. Best practices and the related techniques are more concrete and may change more frequently. Here are some of the common container-related best practices:

- **Aim for small images** - this reduces container size, build time, and networking time when copying container images.
- **Support arbitrary user IDs** - avoid using the sudo command or requiring a specific userid to run your container.
- **Mark important ports** - specifying ports using the EXPOSE command makes it easier for both humans and software to use your image.
- **Use volumes for persistent data** - the data that needs to be preserved after a container is destroyed must be written to a volume.
- **Set image metadata** - Image metadata in the form of tags, labels, and annotations makes your container images more discoverable.
- **Synchronize host and image** - some containerized applications require the container to be synchronized with the host on certain attributes such as time and machine ID.
- **Log to STDOUT and STDERR** - logging to these system streams rather than to a file will ensure container logs are picked up and aggregated properly.

3.7. Kubernetes effect: Benefits



- **Self Service Environments** - enables teams and team members to instantly carve isolated environments from the cluster for CI/CD and experimentation purposes.
- **Dynamically Placed Applications** - allows applications to be placed on the cluster in a predictable manner based on application demands, available resources, and guiding policies.
- **Declarative Service Deployments** - this abstraction encapsulates the upgrade and rollback process of a group of containers and makes executing it a repeatable and automatable activity.
- **Application Resilience** - containers and the management platforms improve the application resiliency in a variety of ways such as:
 - Infinite loops: CPU shares and quotas
 - Memory leaks: OOM yourself
 - Disk hogs: quotas
 - Fork bombs: process limits
 - Circuit breaker, timeout, retry as sidecar
 - Failover and service discovery as sidecar
 - Process bulkheading with containers
 - Hardware bulkheading through the scheduler
 - Auto-scaling & self-healing
- **Service Discovery & Load Balancing & Circuit Breaker** - the platform allows services to discover and consume other services without in application agents. Further, the usage of sidecar containers and tools such as Istio framework allow to completely move the networking related responsibilities outside of the application to the platform level.
- **Declarative Application Topology** - using Kubernetes API objects allow us to describe how our services should be deployed, their dependency on other services and resources prerequisites. And having all this information in an executable format allows us to test the deployment aspects of the application in the early stage of development and treat it as programmable application infrastructure.




4. Amazon EKS and ECS

4.1. Comparison

	EKS	ECS
1 Learning Curve and Community		7
2 Interoperability		
3 Pluggable Architecture		5
Overlay Network (Calico, Flannel, WeaveNet, ...)		
Storage (Ceph, GlusterFS, NFS)		
4 Kubernetes Primitives		
ConfigMap, Secrets, Controller, Pod, Ingress, ...		
5 Security		
L4 Firewalling		
L7 Firewalling		
Network Policy		
Pod Policy		
Service Identity (SPIFFE)		
Throttling		
Certificate Validation		
End to end Monitoring in real time (based on Sidecar)		
6 Service Mesh		
Traffic Management		
Service Identity (SPIFFE)		
Certificate Management		
TLS Termination		
Ingress and Egress		
Observability (Monitoring, Tracing, Logging, Metrics, ...)		
	7	5

5. Delivery on Kubernetes

5. Delivering Distributed Apps: Pods & Services




```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: hello-v1
  namespace: hello
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: hello
        version: v1
    spec:
      serviceAccountName: helloworld-sa
      containers:
        - name: helloworld
          image: istio/examples-helloworld-v1
```

2 Kubernetes Services:

- **ClusterIP**: Exposes the service on a cluster-internal IP. Choosing this value makes the service only reachable from within the cluster. This is the default `ServiceType`.
- **NodePort**: Exposes the service on each Node's IP at a static port (the `NodePort`). A **ClusterIP** service, to which the `NodePort` service will route, is automatically created. You'll be able to contact the `NodePort` service, from outside the cluster, by requesting `<NodeIP>:<NodePort>`.
- **LoadBalancer**: Exposes the service externally using a cloud provider's load balancer. `NodePort` and `ClusterIP` services, to which the external load balancer will route, are automatically created.
- **ExternalName**: Maps the service to the contents of the `externalName` field (e.g. `foo.bar.example.com`), by returning a `CNAME` record with its value. No proxying of any kind is set up. This requires version 1.7 or higher of `kube-dns`.


5. Delivering Distributed Apps: Issues



?

I've deployed my App (Pods & Services) in Kubernetes.
Now, Will the CI/CD scale?,
Is it secure?

6. Traffic Management w/ Ingress



Lab 02: Ingress


Ingress Controller

- 1 Deploy the chosen Ingress Controller in Pods and in a specified Namespace. Default HTTP Backend with custom error pages can be implemented.
- 2 Configure and integrate the deployed Ingress Controller with the EdgeProxy (Firewall, Root Proxy, Load Balancer, etc.). Generally, Ingress Controller is configured as many, during up-time, every API/Microservice will be deployed with Ingress Resource definition (routing). TLS, MTLS or HTTP can be defined.
- 3

```
---  
apiVersion: v1  
kind: Service  
metadata:  
  name: hello-svc-np  
labels:  
  app: hello  
  namespace: hello  
spec:  
  type: NodePort  
  ports:  
    - name: http  
      port: 5030  
      targetPort: 5000  
      selector:  
        app: hello
```

```
---  
apiVersion: extensions/v1beta1  
kind: Ingress  
metadata:  
  name: hello-ing  
  annotations:  
    kubernetes.io/ingress.class: istio  
spec:  
  rules:  
    - http:  
      paths:  
        - path: /hello  
          backend:  
            serviceName: hello-svc-np  
            servicePort: 5030
```

6. Traffic Management w/ Ingress: Issues



?


I've implemented the Ingress Controller in Service Mesh.
Now, Is that Secure?

It isn't secure because the Pods still are accessible.
Any other Pod (container) can call this.


It isn't secure because that doesn't have:

- Separation of Duties
- Principle of Last Privilege
- Identity-based Security
- End-to-end security: the App Container still is available and nobody can stop to make call to that Container.
- Container Patterns (sidecar container as last-mile protection)
- Registry, Auditability and Traceability. It means that all App Containers (APIs and Microservices) living in the Service Mesh should be Identified and Traceable.

7. Traffic Management w/ Ingress + Sidecar



8. Traffic Management w/ Ingress + Sidecar + Network Policy



Lab 04: Ingress, Sidecar and Network Policy

Ingress Controller

- 1 Deploy the chosen Ingress Controller in Pods and in a specified Namespace. Default HTTP Backend with custom error pages can be implemented.
- 2 Configure and integrate the deployed Ingress Controller with the EdgeProxy (Firewall, Root Proxy, Load Balancer, etc.). Generally, Ingress Controller is configured as **Kubernetes Service** or **LoadBalancer**.
- 3 Every API/Microservice will be deployed with Ingress Resource definition (routing). TLS, MTLS or HTTP can be defined.



Secure Traffic Management (L7 Firewall)


- 4 Inject Sidecar in the same Pod where the App Container is living. Sidecar bootstraps security config. to reject traffic from untrusted source.
- 5 The "Single Source of Truth" does validate and authorize the incoming traffic.
- 6 It provides a fine-grained control over all App Containers running in the Service Mesh.
- 7 The TLS termination should be configured properly in order to be managed dynamically for the Ingress Controller. It means to manage the lifecycle of the TLS Certs.



Secure Traffic Management (L3/L4 Firewall)

- 7 The TLS termination should be configured properly in order to be managed dynamically for the Ingress Controller. It means to manage the lifecycle of the TLS Certs.
- 4 Additionally, Sidecar manages the TLS Certificate Lifecycle (revocation, renewal, validations, etc.)

9. Security Assessment – New threat vectors



Kubernetes Security Features

Network policy 1.7

1.7 Stable Using a network plug-in, set and enforce which pods can communicate with each other and other network endpoints. (By default, pods accept traffic from any source.)

RBAC 1.8

1.6 Beta, 1.8 Stable Using RBAC, you can control user and application access to resources. You define a set of roles and their respective permissions, and bind users to roles to grant access. (By default, users start with no permissions.)

PodSecurityPolicy 1.10


1.4 Alpha, 1.10 Beta Using PodSecurityPolicy, define a set of conditions for a pod to be accepted to run in a cluster. These can include run as user, privilege escalation, and SELinux.

1.7 Encrypted secrets

1.7 Alpha Using the EncryptionConfig, you can encrypt data in etcd (including secrets) at the application layer. Choose from AES-CBC, AES-GCM, and secretbox options. (By default, secrets use the 'Identity' EncryptionConfig, which is plaintext.)

1.8 TLS cert rotation

1.7 Alpha, 1.8 Beta For cert authority in your cluster, rotate the kubelet's client and server certs using RotateKubeletClientCertificate and RotateKubeletServerCertificate.



A malicious user with a shell in a container
By default, can very possibly

1. Exfiltrate source code, keys, tokens, and credentials
2. Elevate privileges inside Kubernetes to access all workloads
3. Gain root access to the underlying cluster nodes
4. Compromise other systems and data in the cloud account outside the cluster

Hacking and Hardening Kubernetes Clusters by Example [I] - Brad Geesaman, Symantec

3,025 views

Video Sponsored by redhat

KubeCon CloudNativeCon North America 2017

CNCF [Cloud Native Computing Foundation] Published on 15 Dec 2017


SHARE SUBSCRIBE 9.1K



Lab 05: Security Assessment

10. Service Mesh: Reference Architecture

10. Service Mesh: Istio Framework



Lab 06: Weaving
Service Mesh